

eXtended eXternal Benchmarking eXtension – XXBX

XXBX Getting Started v1.0

Matthew R. Carter

Raghurama R. Velegala
Jens-Peter Kaps
jkaps@gmu.edu

John Pham

April 9, 2018
George Mason University
Fairfax, Virginia



cryptography.gmu.edu



www.gmu.edu

Cryptographic Engineering Research Group

Department of Electrical and Computer Engineering George Mason University
3100 Engineering Building, 4400 University Drive, Fairfax, VA 22030-4444, USA
Voice: (703) 993-1561, Fax: (703) 993-1601

Contents

1	Introduction	3
1.1	Benchmarking	3
1.2	Metrics	3
1.2.1	Throughput	3
1.2.2	ROM	3
1.2.3	RAM	3
1.2.4	Power	3
2	Previous Work	4
2.1	SUPERCOP	4
2.2	XXB	4
3	XXBX System	5
3.1	Overview	5
3.2	XXBX Software – XBS	5
3.3	XXBX Harness – XBH	5
3.4	XXBX Power Shim – XBP	6
3.5	XXBX Device Under Test – XBD	6
4	First Steps	7
4.1	Requirements	7
4.2	Installation	7
4.3	Example Run	7
4.4	Results	9
A	Maybe Something	10

1 Introduction

1.1 Benchmarking

1.2 Metrics

For embedded software cryptography implementations, the areas of interest are ROM usage, Random Access Memory (RAM) usage, throughput, and power. This differs from non-embedded environments, where typically only throughput is considered when evaluating performance, as disk storage (analogous to ROM), RAM, and power are plentiful.

1.2.1 Throughput

Throughput for cryptographic software is given in cycles per byte, in an effort to be independent of different clock speeds within the same family of CPU, as well as variable clock speeds due to power saving [?]. This only helps with variations in clock speeds, or comparing between otherwise identical architectures with varying features such as a Single Instruction Multiple Data (SIMD) unit like NEON or SSE. Results from different instructions per clock and/or instruction formats still cannot be compared with each other.

1.2.2 ROM

ROM usage is typically determined by the program code as well as any constants, and preinitialized variable data. Microcontrollers have anywhere between a few kibibytes to a mebibyte of ROM. This typically consists of the `.text` (code and constants) and `.data` (initialization data for static buffers in RAM) sections in a compiled binary.

1.2.3 RAM

RAM is organized into statically allocated RAM set aside at the beginning of program execution, stack memory that grows as function call chains get deeper, and dynamic allocations via `malloc` and similar functions. The statically allocated RAM is divided into `.bss` and `.data` sections. `.data` is for preinitialized modifiable buffers in RAM that also take up space in ROM for the initialization data as previously mentioned, while `.bss` is for uninitialized buffers that only occupy RAM.

1.2.4 Power

As software ultimately runs on hardware containing a processor, the same constraints on power that affect hardware affects software. The power and energy used is influenced by which processing units the software uses, the microprocessor the software runs on, the clock speed and the total running time.

2 Previous Work

There are many existing tools for benchmarking cryptographic software. The System for Unified Performance Evaluation Related to Cryptographic Operations and Primitives (SUPERCOP) is a cryptographic benchmarking tool that benchmarks a large variety of cryptographic primitives on general purpose computers[?]. However, it is not capable of being run on embedded platforms, and lacks important metrics for embedded such as ROM and RAM usage.

Two software benchmarking tools for embedded are XBX[?] and FELICS [?][?]. XBX is an tool for evaluating hash functions packaged for SUPERCOP, and primarily runs implementations on actual hardware.

FELICS is a newer tool, focused on lightweight primitives, and current lightweight block ciphers, and adds cycle-accurate simulation for a variety of platforms as well as running implementations on actual hardware.

In this paper, we further extend XBX to cover AEAD primitives in addition to XBX, and allow easy extensibility to support the rest of the operations SUPERCOP supports. In addition, we intend to add the ability to measure power and energy consumption.

2.1 SUPERCOP

SUPERCOP “[...] is a toolkit developed by the VAMPIRE lab for measuring the performance of cryptographic software [?].” The project collects and benchmarks many implementations across many primitives on multiple platforms. It consists of a series of shell scripts and test harnesses that compile implementations and their dependencies across different compilers and flags. It then verifies outputs and correctness, and measures performance characteristics. Power measurements, code size, and memory usage data measurements are not supported.

2.2 XBX

The eXternal Benchmarking eXtension (XBX)[?] is an extension to SUPERCOP that covers platforms that are not capable of hosting a compiler and require cross-compilation, such as microcontrollers or small embedded Linux systems. It reuses the primitive implementations collected by the SUPERCOP project (which it refers to as an “algopack”), as well as others, and attempts to retain the same output data format. XBX is able to test multiple implementations across multiple primitives with multiple compilers and compiler flags without intervention, which is important with the large number of combinations to be tested. Due to the importance of small code and memory size for embedded platforms, these are also measured and recorded.

3 XXBX System

3.1 Overview

The XXBX system is comprised of four main components:

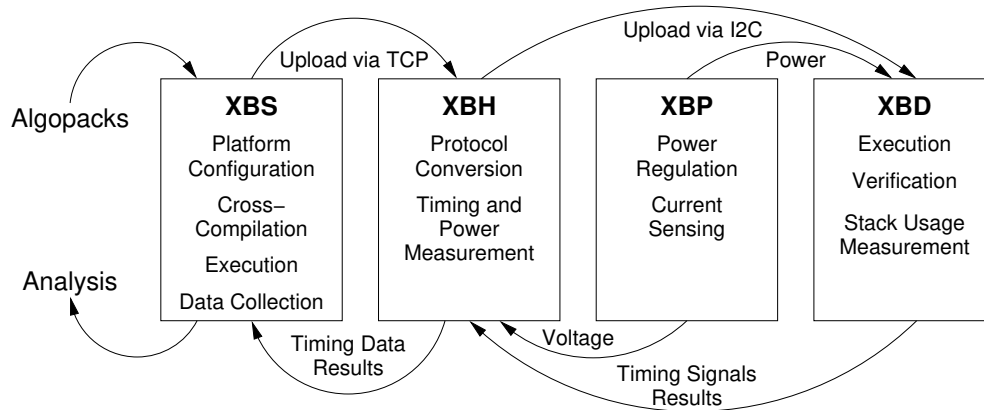


Figure 3.1: Block Diagram of the XXBX System

The communication between the XBS and XBH uses TCP/IP over Ethernet, while I2C is used between XBH and XBD. A buffer is used to queue the commands received from the XBS before sending them to XBD.

3.2 XXBX Software – XBS

The XBS is software running on a PC that consists of scripts that handle compilation of algo packs and firmware upload to the XBH, as well as data aggregation and logging. Static allocation of RAM and ROM usage is also calculated from the compiler output. Multiple compilers and compiler flags are used to compile an implementation.

3.3 XXBX Harness – XBH

The XBH is a board that translates high level commands from the XBS sent via TCP/UDP to series of commands over a protocol understood by the XBD (I²C , RS232, or Ethernet). Future references to the XBS communicating to the XBD refers to this. It also measures execution time via the timer capture pins, which precisely record the time that an edge event occurs. Because not all XBD platforms support cycle counters, the execution time is measured externally.

3.4 XXBX Power Shim – XBP

The power is measured using a shunt resistor circuit. The current is amplified here so that ADC on the XBH can measure it.

3.5 XXBX Device Under Test – XBD

The XBD is the platform hardware that the primitive implementations are tested on. It performs the execution task and notifies the XBH of start and completion via a general-purpose input/output (GPIO) pin connected a timer capture pin on the XBH. It also returns the maximum stack RAM used to the XBH over the XBD↔XBH communication channel. The supported platforms are ARM, AVR, MSP430, and MIPS.

The XBD code consists of two components. One is a small bootloader that is flashed once to XBD that performs basic calibration and self programming of the primitive implementations over the communications channel.

The second component is a test harness that is linked with the primitive implementation that is downloaded over the communication channel. This code is responsible for signaling start and stop, supplying the implementation with input data, verifying the correctness of the algorithm implementation and returning the output to the XBH.

Both of these are linked to code which provides device specific drivers for communication and execution signaling, as well as stack measurement.

4 First Steps

4.1 Requirements

4.2 Installation

4.3 Example Run

Prior to running any scripts, the XBD is initialized by compiling the platform-specific bootloader and flashing it normally (i.e. using a programming cable). This bootloader has the responsibility of downloading code from the XBH. This allows board-specific programming tools to be avoided, besides the initial setup, and thus all subsequent operations can be performed by the XBS over a network, potentially as part of a farm.

Scripts from the XBS compile all implementations of primitives, linking with the hardware abstraction layer (HAL) and the test harness, with the possible {compiler, flags} defined in the platform support files. The number of generated application binaries should be {compilers, flags} \times primitive implementations. The UNIX `size` command is run on the generated application binaries to get the sizes of `.bss`, `.data`, and `.text`. Implementations listed in a platform-specific blacklist and a global blacklist are omitted from compilation. Binaries are compiled into Executable and Linkable Format (ELF) which is then converted into Intel Hex (IHEX) for transfer to the XBH. XBH checks to see if a binary will fit, otherwise it will not be loaded later.

The XBS then orders the XBH to calibrate cycle calculations on the XBD. The XBD's bootloader executes a busy loop for a fixed number of cycles determined by the HAL. The XBH measures the physical time executed. The cycles run is then calculated by dividing the measured time by the clock rate.

The XBS then sends the application binary to the XBD. This is sent one flash block at a time, as existing data on the flash must be erased, and this can only be done one block at a time. After each block is completely received, the XBD writes the block from the RAM buffer to flash memory.

When all blocks are received, the XBD will switch to the application by calling the address of the application's entry point, which is set to a fixed address by linker scripts or compiler flags. This is the code that copies `.data` from ROM to RAM, zeroes out `.bss`, and initializes the stack pointer, then calls `main`. This is typically defined by the C runtime, and is often known as "crt0". However the ARM Stellaris platform has this defined as part of the HAL, as it is based off example code provided by the chip vendor which also does this. As the initialization is run, the previous call stack and static memory allocations of the bootloader is obliterated and replaced with that of the application. A similar procedure is applied to switch back to the bootloader after application execution.

Once in application mode, the XBD is ordered to verify correct execution of the primitive implementation. This is done by hashing repeatedly (XBH only supports hashing), mixing in the results of the previous hash, and comparing the results to that generated by the reference implementation on a PC.

The XBD is then ordered to perform the actual benchmarks. Data randomly generated by the XBS of varying lengths (0, 1, 2, 4, 8, 16, 32, 64, 128, 256, 512, 576, 1024, 1536, and 2048 bytes) is sent to the XBD, and hashed. The XBD signals to the XBH the start and stop times by setting a signal line low upon start and raising it high again upon stop. Any delays in detecting timing should be symmetrical across start and stop and should not affect the result [?]. This is repeated multiple times (defaulting to 5).

Stack usage is measured by painting memory with a canary value from the top of the statically allocated memory sections to the current location of the stack before executing the implementation. After execution, the number of addresses still containing the canary value and not overwritten by the call stack growing when the implementation returns is subtracted from the number of addresses painted to determine the amount of stack memory used.

Upon execution completion, the XBS will request the measured stack value from the XBD and the timing results from the XBH. These values are logged and the process repeats starting with loading the next application binary, until all of them are tested.

Results are output

First, the configuration in `config.ini` is edited. This specifies whitelists and blacklists of implementations, the target operation, the target platform, paths, parameters, dependencies, etc. We assume that resuming a failed run will not involve changing the target operation and platform. Logging levels can be configured in `logging.ini`, and uses Python's built-in logging framework.

Then `compile.py` is then run. This creates the file `data.db`, used to store results, configuration, and other detailed information. `config.ini`, a platform-specific config file, and various blacklist files are all parsed.

If a whitelist exists in `config.ini`, only the implementations in this list will be built (and later run), otherwise all implementations available in the specified operation and primitives are built, excluding those specified in the blacklist. There is a blacklist in `config.ini`, specific to the run, as well as a global blacklist for implementations that are broken, and platform specific blacklists for implementations that are broken on specific platforms. Hashes of all implementation and platform code are taken to ensure exact replication of results, if required.

Once the configuration is assembled, the compile script generates the makefiles and header files for each implementation and builds the code for the XBD for each implementation \times {compiler, flag}. This proceeds quickly, as the script spawns as many subprocesses as there are cores to execute builds. We decided to run multiple makefiles simultaneously as opposed to running GNU make with the `-j` flag, which does parallelization at the makefile level, as even with `-j`, the makefiles would stall trying to build something that would not build, e.g. attempting to compile x86 assembly to ARM. Running separate makefiles in parallel builds more quickly, by allowing stalls to not block other tasks. The `size` command is run as in original XBX on the compiled binaries, and the values saved in the database. We decided to limit the maximum static memory allocations to 3/4ths the total available RAM, leaving the remainder for stack growth, for both the Tiva-C and the MSP430 XBDs in the platform configuration.

As the builds use GNU `make`, rebuilds do not require recompiling everything from scratch, only the code that changed or that was interrupted. Original XBX did not use makefiles. It copied all source files to a build directory, and compiled the source files together all at once, which meant all files are always rebuilt even if nothing changed.

Unlike XBX, and like SUPERCOP, primitive implementations can depend on other implementations. However, the relations have to be explicitly defined, as there are multiple metrics for what

a good implementation is, unlike SUPERCOP.

Once a build succeeds, the `execute.py` script is run. The configuration used from the last build is loaded from `data.db`, if `config.ini` has not changed between the last build and calling `execute.py`. Behavior is undefined if `config.ini` or blacklists or code are modified between calls of `compile.py` and `execute.py`. Calibration is run as in original XBX.

For each successfully compiled implementation, the binary is loaded into the XBD and tests are run, as mentioned in section ???. These tests can be run multiple times if configured for such, however they take a considerable amount of time. Afterwards, the XBS generates data to send to the XBD for performance testing, as specified in the config file. A python object defining the parameters, how to generate them, and how to handle what is returned by the XBD is required to support an operation, and as previously mentioned, we have hash functions and AEAD implemented.

In the case of hashing, we simply generate random data of the specified lengths and send them to the XBD, discarding the result, while timing and measuring execution, as we don't have a known answer for hashing random data, while for AEAD, we encrypt the data, get the ciphertext from the XBD, and send it back for decryption, as well as decryption of a tampered ciphertext. We log time (and later power) measurements for all these results. We do an additional verification to see if the decryption matches the unencrypted plaintext, as this is essentially free without additional computations as the XBS has this information available. All of this information is logged for later analysis.

4.4 Results

XXBX results can be analyzed by opening the generated SQLite database (`data.db`) in a SQLite browser application, such as DB Browser for SQLite [8]. For analysis, the SQLAlchemy objects can be manipulated directly in an IPython [11] notebook. IPython is an interactive python environment, with an interface similar to Maple or SageMath with notebooks. Information was saved into a SQLite [12] database.

A Maybe Something